

Archive Search — Quick Question: Flagging Old Owner Instructions as Discriminatory

Prepared for: Peter McKenzie, Rincon Management, to forward to outside counsel

Date: 2026-09-13

Purpose: one specific new question, separate from the privilege question above and from any prior approval of live-mail handling.

What's already approved, for context: for a *new* owner email that gives Rincon a discriminatory instruction (e.g., "don't rent to voucher tenants"), counsel already approved a fixed response: Rincon's system logs the instruction and a note stating Rincon can't comply with it and will follow standard, non-discriminatory procedure. That's for something happening right now, where the note documents a real, current refusal.

What's new: we want to run the same check backward, against old email — some of it years old. When the AI finds an old instruction that looks discriminatory, it would write and permanently save a note today, saying something like: "AI-assessed in 2026: this 2022 instruction, if acted on, would require Rincon's standard refusal — no record confirms what was actually communicated to the owner at the time." Unlike the live version, this note doesn't describe anything Rincon actually did — it's a new, present-day AI opinion about a named owner's old email, saved into Rincon's own records, with nobody having reviewed it first.

The question: is it safe for Rincon's system to generate and save that kind of note automatically, unreviewed, about old mail — or should the AI only flag the category and send it to a human, with a person deciding whether any note gets written at all? We want your read specifically because if the AI misjudges an old, ambiguous instruction, Rincon's own database ends up containing an AI-generated accusation about a real owner that nobody checked before it was saved.

Prepared with AI assistance (Claude / "Jarvis") for Rincon Management. Not legal advice — prepared to be reviewed by qualified California outside counsel.