

Archive Search — Rollout Checklist

Purpose: one single, plain-English sequence for this entire project, so nothing gets lost across a long session or a gap between sessions. Each step links back to the real document that has the actual detail — this file is the map, not the map's territory. Update the checkbox and date whenever a step actually happens; don't let this drift out of sync with reality.

The three source documents this checklist summarizes:

- `projects/hub/email-intake/archive-search-v1-scope.md` — the product design (what this is, who can use it)
- `projects/hub/email-intake/archive-search-technical-spec.md` — the engineering design (schema, screening logic, safety mechanism)
- `compliance/archive-search-ai-risk-assessment.md` — the formal risk document

The Sequence

- ☒ **1. Product scope written** — 2026-09-10
- ☒ **2. Product scope reviewed by Asimov and Mason** — 2026-09-10. Real findings, all resolved (hold-policy conflict, access-model correction, keyword gaps identified).
- ☒ **3. Technical spec written, addressing every review finding** — 2026-09-10
- ☒ **4. Technical spec reviewed by Neo** (schema-level) — 2026-09-10. Three real problems found and fixed (enforcement gate, retention-clock gap, table-locking risk).
- ☒ **5. Technical spec re-reviewed by Asimov and Mason** — 2026-09-10. Approved with conditions; both conditions (the CI guardrail requirement, the risk-assessment correction) already fixed in the documents.
- ☐ **6. Schema + screening-pass code built** (Neo, Q, TARS, Judge) — **in progress now**. Does not touch real data. Not the search UI.
- ☐ **7. Peter applies the real database migration** in Supabase's SQL Editor — waiting on step 6 to finish and get a clean Judge verdict.
- ☐ **8. Peter explicitly authorizes running the screening pass against real data** — a separate decision from step 7. Nothing runs automatically. This is the step that actually touches all 254,000+ real messages for the first time.
- ☐ **9. Validation sample reviewed** — a real person reads ~1,000 real messages (the stratified sample the technical spec defines) to confirm the screening pass didn't miss anything. **Zero-miss exit rule:** even one real miss means fixing the root cause and re-running the *entire* pass, not just patching around it.
- ☐ **10. Held-bucket review** — whatever the pass found as "held" (real legal/Fair Housing matters) gets reviewed once by Peter, the DO, and possibly outside counsel, separate from search itself.
- ☐ **11. The actual search tool gets designed and built** — the search box, the results screen. Not started. Waits on step 9 clearing.

- ☐ **12. Initial access: narrow** — limited to whoever did the step 9 validation review.
- ☐ **13. Access widens to the full 8 confirmed people** (7 matched Missive members + Caylee as the accepted exception) — once step 6-9's mechanisms have run cleanly for a stretch, not on day one.
- ☐ **14. (Recommended, not blocking) Ongoing re-validation** — a periodic, smaller check on the ongoing screening of new mail once this is live, mirroring step 9 at a lighter scale (Mason's recommendation).
- ☐ **Standing, not a one-time step:** the real human-issued litigation-hold obligation (independent of this whole system) stays in force — if Rincon has actual notice of or reasonably anticipates litigation touching specific tenants/units, a human-issued hold is required regardless of what the automated screening tagged.

What This Checklist Is Not

Not a new design document. Not a substitute for reading the three real source documents when the actual detail matters. Just a durable, single place to answer "where are we" without re-deriving it from a long conversation.