

Complaint & Issue Tracking — V1 Scope

Prepared for: Stephen's review, alongside Peter's own thinking

Date: 2026-09-08

Status: Draft scope for discussion — not yet built, not yet formally approved. Written to capture a full brainstorming session between Peter and Claude; the goal now is to get a second set of eyes and sharpen it further before it becomes a real build plan.

One assumption baked into this document, stated plainly: everything below that involves reading email content (matching a message to a property/tenant/owner by content, and the significance tagging itself) still needs to clear a Fair Housing legal/governance review before it can actually go live — that review hasn't happened yet. For the purposes of this scope document, we're assuming it clears, so the conversation with Stephen can focus on whether this is the *right tool*, not get stuck on a compliance step that's a separate, parallel track. If that review comes back with real restrictions, this document may need to change.

1. The Problem

Rincon has long struggled to track complaints effectively. Today, a complaint exists as scattered messages in someone's inbox or a phone call nobody wrote down — there's no single record of it, no clear owner, and no way to tell whether it's actually been resolved or just gone quiet. Some of what's been missing is *noticing* a complaint at all; some of it is *nobody owning the follow-through* once it's noticed.

Separately, Rincon has been backfilling years of historical email (Faria and Solimar Team Inboxes in Missive) into a searchable archive. That project is a source of real value here — not just going forward, but retroactively, since years of real correspondence already exist and can be mined once the tools below are built.

2. The Core Design Principle

Not everything that comes in is a big deal, and burying the tool in routine noise defeats the purpose. The explicit goal is: track every complaint for the record, but only prominently surface the ones that are actually significant — a "big deal," in Peter's words. An AI system can narrow down what a human needs to look at, but it should never unilaterally decide something isn't worth seeing at all. When it's genuinely unsure, it should say so and let a person make the call, not guess confidently in either direction.

The mechanism for "unsure," confirmed with Peter: uncertainty gets its own flag — "Needs a human call" — separate from the six real categories in Section 3. It routes to the exact same place a confirmed big-deal item goes: the DO's queue, since the DO already owns big-deal triage. It's visibly marked as uncertain rather than silently dropped into a category it might not belong in, and it's never sent to a new, separate review queue nobody's already watching.

Fair Housing detection, confirmed with Peter: reuses Maintenance History's real content-screening tool — the more advanced of the two Fair Housing screens in this codebase, not the simpler keyword-only scanner originally

proposed. It's a genuine two-layer system (a keyword list, plus a second, independent AI layer that can catch something no keyword matches) that just went through a real precision redesign with Mason and outside counsel (2026-09-05): six specific words that caused most of its false alarms (e.g. "white"/"black" almost always meaning a paint color or a water-heater brand, not race) now get one targeted AI disambiguation question before flagging, cutting false positives sharply while still defaulting to flag-it whenever genuinely unclear. **Behavior, deliberately different from Maintenance History's own "held until reviewed" gate:** a Fair-Housing-flagged complaint stays fully visible in the normal complaint view, tagged, same "always visible" principle as everything else in this doc — it routes into the same "Needs a human call" DO queue described above rather than being withheld from ordinary view. Peter's explicit call: the system never blocks or hides anything; a human always decides what happens next, but they decide with the complaint already in front of them, not gated behind a separate release step.

Nothing important should ever be buried behind a click-through or a dropdown menu. If something is significant enough to surface, it should be visible the moment the relevant record loads — flat and prominent, not hidden behind an extra click.

Sharpened after seeing the first mockup, then refined again on where each rule applies: Property 360's Issues section (and the Hub home-page count) shows big-deal items only — full stop, not even collapsed behind a toggle. That's a property-record page someone opens for a fast, reliable read on a property; routine noise doesn't belong there in any form. The **central complaints tool is different:** big-deal items are still surfaced prominently up top, but routine, non-big-deal items are also visible there — tucked behind a dropdown/expand, a deliberate look rather than something that shows up passively, but not hidden from the tool entirely the way they're hidden from Property 360. Routine items are tracked in the underlying system either way (Section 3's "track everything" still holds — recurrence detection needs the full record regardless of where anything is displayed).

3. What Counts as a "Big Deal" — The Surfacing Categories

Every complaint gets tracked regardless of size. Only these get prominently flagged as needing attention:

1. **Legal / compliance exposure** — attorney involvement, code enforcement, a habitability or Fair Housing complaint, anything that could become real liability if ignored.
2. **Blocked resolution** — this is not "a repair is needed," which is just routine work Rincon handles constantly. It's specifically when the normal path to fixing something breaks down: an owner gives a hard refusal to authorize or pay for something, a tenant refuses to allow access for a repair, or — importantly — **2 or more days of silence** (shortened from an initial 4-day proposal — confirmed with Peter that 4 days was too long a window before something counts as blocked) after Rincon has asked for authorization or access counts the same as a refusal, even without an explicit "no." This runs on its own clock, separate from the aging/stall clock in Section 4 — see that section for the split.
3. **Relationship / churn risk** — an owner hinting they're unhappy with management, mentioning other companies or selling; a tenant threatening to break the lease or withhold rent.
4. **Escalation / recurrence** — the same issue coming up again, even if each individual instance looked routine. This is the single best predictor of a real problem, and it requires the full tracked record (including "resolved" items) to detect at all.

5. **Major money / property risk** — not routine spend, but something at the scale of a roof replacement or an insurance-claim-level event.

6. **A one-off instruction from an owner that falls outside normal procedure** — e.g., "always call me before any repair over \$200," "don't ever contact this tenant directly." These matter because staff who weren't there when the instruction came in will default to standard process and unknowingly break it — and because these are exactly the kind of statement that occasionally carries real Fair Housing risk hiding inside a casual instruction. Detecting one of these should not just flag the email — it should propose a new entry in Rincon's existing Owner & Tenant Operational Notes system, since that's the durable, reviewed place this kind of standing rule actually belongs, not an email thread nobody will think to search months later.

Explicitly not flagged, on purpose: routine maintenance that gets handled without drama, scheduling coordination, payment confirmations, standard renewal paperwork, ordinary vendor coordination — even though there's a lot of it.

To be explicit for review purposes: an AI model reads the full content of tenant and owner correspondence to assign these categories. No email is pre-filtered by a human before the AI sees it — the AI's output is what a human reviews, not a gate on what the AI is allowed to read.

3a. Frustration / Tone — A Fourth Tagging Dimension

Added after further discussion with Peter: **something that starts out small can turn into a bigger issue, and the earliest signal of that is often tone, not content.** A tenant's third message about the same small thing, worded more curtly than the first two, is worth surfacing even before it technically qualifies as any of the six categories above.

Unlike the "predictive" capabilities in Section 8 that need months of historical data to establish a baseline, this doesn't need one — it's a read of a single conversation's own message history, not a comparison against other conversations. That makes it realistic for v1, not a later fast-follow: alongside resolution status, pattern, and category (Section 3.2 of the original design discussion), add a fourth dimension — a frustration/tone signal, computed by having the AI read the conversation's messages in order and note whether things are getting more strained over time, not just what the current message says in isolation. Same posture as everything else here: it narrows what a human should look at, it never silently decides on its own that something's fine.

4. Structure — Turning a Complaint Into a Real, Trackable Thing

- **A real lifecycle, with an accurate history — not just a current-status field.** Open → In Progress → Blocked → Resolved. "Resolved" requires an actual note on what was done — not just disappearing from the list. Confirmed with Peter: every stage change gets logged — timestamp, who moved it, and (for anything other than a straight Open → In Progress) a short note on why, e.g. what's blocking it or what actually got done at Resolved. That's what makes the record trustworthy all the way to resolution: the DO or anyone reviewing later can see the real path a complaint took, not just trust whatever the current status field says.

- **An owner from the moment it's created — and for big-deal items, that owner is the DO by design, not just by default.** Refined after discussion with Peter: the DO owns the *process*, not necessarily the fix. When a big-deal complaint lands, the DO researches it and decides the right course of action using their own judgment

— that might mean resolving it themselves, or it might mean delegating to the right person at the pod level. Either way, the DO stays accountable for it; delegating doesn't hand off ownership, it just changes who's doing the work.

- **An aging clock, not just a snapshot — and notification is simpler than a multi-tier ladder because of the point above.** If a big-deal item stalls, the DO gets notified — whether they're sitting on it personally or it's stalled with whoever they delegated it to. One tier (the DO), not an escalation chain climbing up through multiple people. Exact timing (how long before the first nudge) still open. **Two separate clocks, confirmed with Peter, doing two different jobs:** the 2-day silence clock in Section 3's "Blocked resolution" is a *classification* trigger — it decides whether a complaint counts as blocked at all. This aging clock is an *escalation* trigger — it only starts once something is already a confirmed big-deal item, and measures how long it's sat before the DO gets nudged. A complaint can trip the first clock without ever needing the second (getting classified as blocked doesn't necessarily mean it then sits unaddressed), and the second clock only exists for items that already cleared the big-deal bar.

- **Every complaint has a subject** — a property, a tenant, an owner, *or a team member*. Not every complaint is about a property (a complaint about how a specific staff member handled something is real and worth tracking, and isn't tied to a property at all). **Owner-tied and tenant-tied items must be visually separated wherever they're shown together** — obvious at a glance, not something staff have to read through each item to decipher. Layout approach pending Tron's design suggestion.

- **Multiple intake paths.** Complaints will keep coming in by more than just email — phone (Aircall is already connected), in-person conversations, possibly AppFolio's own tenant messaging someday. The record needs to support all of these, with the email-detection work as one automatic source among what should eventually be several, plus a fast, simple manual-entry path for anything else. If logging something manually takes more than a few seconds, staff won't use it, and the "one central place" idea fails in practice.

- **Automatic cleanup, never deletion.** Once a tenant's lease ends (already a real, synced fact from AppFolio), their tracked items should drop out of the active view automatically — kept on record, never shown as something needing action.

- **No match found still creates a record, confirmed with Peter.** A complaint that can't be tied to any known property, tenant, owner, or vendor is still created — as an orphan record, visible in the central tool, flagged "needs matching" — rather than silently failing or being dropped. A person can attach it to the right property/person by hand later. Same "never silently drop anything" principle the rest of this doc already applies everywhere else.

- **A vendor field, when one's involved — linked to the real vendor list, not free text.** Distinct from "subject" (who the complaint is *about*): this is which vendor, if any, was involved in what happened. Links to Rincon's existing AppFolio vendor list, the same reasoning as tenant/owner matching — "ABC Plumbing" typed three different ways would fragment pattern detection if it were just text. Populated two ways: a simple dropdown/autocomplete for manual entries, or the same content-extraction approach used to identify a property/person in an email, pointed at vendor names instead. Left null when no vendor is identifiable — same clean-omit rule as everything else here. This is the one genuinely new field in this design with no existing data to lean on (contrast with team-member pattern detection below, which reuses data that already exists).

5. Where This Lives

- **A new, dedicated tool in the Hub** is the canonical home for every complaint, regardless of source or subject — the place to go for "show me everything," filterable by property, team member, category, status.

- **Contextual surfacing on top of that** — a property-tied complaint also shows up on that property's own Property 360 page. **Team-member-tied complaints, confirmed with Peter: no dedicated profile page for v1.** Rincon doesn't have one today, and building one is a bigger, separate feature — the central tool's own filter-by-team-member view is enough to start. A dedicated page can be a fast-follow if it turns out people actually want it, not something built speculatively now.
- **A live count on the Hub's home page** — the first thing anyone sees regardless of which tool they're actually heading to, so a real, open issue can't just sit unnoticed because nobody thought to check the dedicated tool.

6. The Director of Operations Connection

The DO's core metric is churn rate, and complaints are a direct input into churn. That makes the DO the natural default owner for the big-deal categories — particularly churn-risk and blocked-resolution — not just one more person in a shared queue. A simple rollup view (how many open, how many overdue, by whom) gives the DO's churn story something concrete to point to, instead of a gut feeling.

7. Relationship to the Existing "Owner in Distress" Process (LeadSimple)

This already exists and is genuinely active — 80 real cases, the large majority already resolved, tracking real fields like "Did Owner Approve Work?" (Yes/No/Undecided) and "What is the reason for distress?" (Maintenance, Vacancy, Customer Service, Owner Mentioned Selling).

Important distinction, confirmed with Peter directly: this new complaint system is *not* a front door into Owner in Distress, and most complaints should never touch it. Owner in Distress is one specific **escalation destination** for the subset of complaints that genuinely indicate real owner distress (most naturally, the churn-risk category) — not the general home for every tracked complaint. Most complaints should be handled and closed entirely within the new system.

The escalation mapping, confirmed with Peter: only **relationship/churn risk** (Section 3, category 3) escalates to Owner in Distress — it's the one category that actually matches what Owner in Distress already tracks (maintenance, vacancy, customer service, owner mentioned selling). The other five categories (legal/compliance, blocked resolution, escalation/recurrence, major money/property risk, one-off owner instructions) stay entirely within the new system — they're Rincon-process issues, not owner-distress signals, and don't route anywhere else. **Who decides:** the DO — consistent with them already owning both this system's big-deal triage and the churn metric Owner in Distress ultimately feeds.

8. The World-Class Vision — All of It, Sequenced

Peter's explicit direction: build toward all of the following, don't cut any of them — but sequence them honestly rather than trying to build the hardest pieces before there's any real data to make them trustworthy.

In v1 — these are workflow and design decisions on what's already scoped above, not new hard problems:

- The structured lifecycle, ownership, and aging/escalation clock (Section 4)
- Fast, frictionless manual intake
- Comfort with uncertainty — ambiguous cases go to a human, never a confident wrong guess
- A simple visibility rollup and the home-page count (Sections 5–6)

Fast-follow, likely weeks rather than months after v1 ships — because the historical email archive means this doesn't have to wait for new data to accumulate:

- Recurrence and pattern detection ("has this come up before") — can run retroactively against years of already-archived correspondence once matching and tagging are live
- Root-cause/pattern detection across three axes, confirmed in scope after discussion with Peter — Rincon is a small enough team that connecting patterns to specific people makes sense, not something to shy away from:
- **A problem vendor** — the same vendor's jobs generating complaints across multiple, unrelated properties (needs the new vendor field above).
- **A problem property** — one property generating more tracked complaints than similar properties, regardless of who manages it.
- **A coaching signal** — one team member's properties showing a disproportionate share of escalations, using the assignee field already in Section 4 (no new field needed), or by joining against whoever Rincon already has on record as each property's manager (also no new field — reuses existing data rather than duplicating it).

Needs real, live operating time before it can be trusted — genuinely can't be rushed:

- Operational effectiveness metrics (actual resolution speed, how well the DO is closing things out) — needs the new lifecycle actually in use for a few months, since the old emails were never tracked against these concepts
- True predictive-of-churn modeling (catching distress before it's explicit) — needs enough real churn events to have happened to validate any pattern, and those are inherently slow to accumulate; realistically 6–12 months of real operating history

Also part of the vision, not yet scoped in detail:

- Closing the loop with the tenant or owner who complained (a status update along the way, cutting down on frustrated follow-up complaints) — doesn't need automation to start; even a manual habit changes the relationship. **Confirmed with Peter: stays as-is, out of v1, parked for later** — real outbound communication, gets its own separate review when it's actually built.
- Searchability across the full email archive — technically buildable independent of everything else here, but still behind the same Fair Housing gate before real use

9. Open Questions for Stephen

1. Does the big-deal category list (Section 3) match how you'd actually describe a real "big deal" from where you sit, or is something missing or overweighted?

2. ~~Where's the actual line between "tenant is slow to respond" and "tenant is refusing" for the blocked resolution rule — is the timing right, or does it depend on the type of issue?~~ RESOLVED — Peter set this at 2 days (down from an initial 4-day proposal). Still worth Stephen's read on whether 2 days is right across every issue type, or should vary by severity.
3. Which categories should escalate to Owner in Distress specifically, versus a different path entirely (Section 7)?
4. Does Rincon need a team-member profile/record page for team-member-tied complaints to surface against, or is the central tool enough for that case (Section 5)?
5. Anything from day-to-day complaint handling that this scope is missing entirely?